Machine Learning
(Học máy – IT3190E)

Khoat Than

School of Information and Communication Technology

Hanoi University of Science and Technology

2022

Contents

- Introduction to Machine Learning
- **Supervised learning**
    - **Support Vector Machines**
- Unsupervised learning
- Reinforcement learning
- Practical advice

Support Vector Machines (1)

- **Support Vector Machines (SVM)** (máy vectơ hỗ trợ) was proposed by Vapnik and his colleages in 1970s. Then it became famous and popular in 1990s.

- Originally, SVM is a method for linear classification. It finds a hyperplane (also called *linear classifier*) to separate the two classes of data.

- For *non-linear classification* for which no hyperplane separates well the data, *kernel functions* (hàm nhân) will be used.

  □ Kernel functions transform the data into another space, in which the classes are linearly separable.

- Sometimes, we call linear SVM when no kernel function is used. (in fact, linear SVM uses a linear kernel)

Support Vector Machines (2)

- SVM has a strong theory that supports its performance.

- It can work well with very high dimensional problems.

- It is now one of the most popular and strong methods.

- For text categorization, linear SVM performs very well.

1. SVM: the linearly separable case

- Problem representation:
  - Training data $\mathbf{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), …, (\mathbf{x}_r, y_r)\}$ with r instances.
  - Each $\mathbf{x}_i$ is a vector in an n-dimensional space, e.g., $\mathbf{x}_i = (x_{i1}, x_{i2}, …, x_{in})^T$. Each dimension represents an attribute.
  - Bold characters denote vectors.
  - y_i is a class label in $\{-1; 1\}$. '1' is possitive class, '-1' is negative class.
- **Linear separability assumption:** there exists a hyperplane (of linear form) that well separates the two classes (giả thuyết tồn tại một siêu phẳng mà phân tách 2 lớp được)

Linear SVM

- SVM finds a hyperplane of the form:

$$f(\mathbf{x}) = \langle \mathbf{w} \cdot \mathbf{x} \rangle + b \qquad \text{[Eq.1]}$$

  - $\mathbf{w}$ is the weight vector; b is a real number (bias).
  - $\langle \mathbf{w} \cdot \mathbf{x} \rangle$ and $\langle \mathbf{w}, \mathbf{x} \rangle$ denote the inner product of two vectors (tích vô hướng của hai véctơ)

- Such that for each x_i:

$$y_i = \begin{cases} 1 & if \quad \langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b \geq 0 \\ -1 & if \quad \langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b < 0 \end{cases} \qquad \text{[Eq.2]}$$

Separating hyperplane

- The hyperplane (H_0) which separates the possitive from negative class is of the form:

$$\langle \mathbf{w} \cdot \mathbf{x} \rangle + b = 0$$

- It is also known as the *decision boundary/surface.*

- But there might be infinitely many separating hyperplanes. Which one should we choose?

[Liu, 2006]

Hyperplane with max margin

- SVM selects the hyperplane with **max margin**. (SVM tìm siêu phẳng tách mà có lề lớn nhất)

- It is proven that *the max-margin hyperplane has minimal expected errors among all possible hyperplanes.*

[Liu, 2006]

Marginal hyperplanes

- Assume that the two classes in our data can be separated clearly by a hyperplane.

- Denote ($\mathbf{x^+}$,1) in possitive class and ($\mathbf{x^-}$,-1) in negative class which are *closest* to the separating hyperplane H_0 ($\langle \mathbf{w} \cdot \mathbf{x} \rangle + b = 0$)

- We define two parallel *marginal hyperplanes* as follows:

  □ H_+ crosses $\mathbf{x^+}$ and is parallel with H_0: $\langle \mathbf{w} \cdot \mathbf{x^+} \rangle + b = 1$

  □ H_- crosses $\mathbf{x^-}$ and is parallel with H_0: $\langle \mathbf{w} \cdot \mathbf{x^-} \rangle + b = -1$

  □ No data point lies between these two marginal hyperplanes. And:

$$\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b \geq 1, \qquad \text{if } y_i = 1$$

$$\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b \leq -1, \qquad \text{if } y_i = -1 \qquad \text{[Eq.3]}$$

The margin (1)

- **Margin** (mức lề) is defined as the distance between the two marginal hyperplanes.

  □ Denote d_+ the distance from H_0 to H_+.

  □ Denote d_- the distance from H_0 to H_-.

  □ ($d_+ + d_-$) is the margin.

- Remember that the distance from a point $\mathbf{x_i}$ to the hyperplane H_0 ($\langle \mathbf{w} \cdot \mathbf{x} \rangle + b = 0$) is computed as:

$$\frac{|\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b|}{\|\mathbf{w}\|} \qquad \text{[Eq.4]}$$

  □ Where:

$$\|\mathbf{w}\| = \sqrt{\langle \mathbf{w} \cdot \mathbf{w} \rangle} = \sqrt{w_1^2 + w_2^2 + \cdots + w_n^2} \qquad \text{[Eq.5]}$$

The margin (2)

- So the distance d_+ from $\mathbf{x}^+$ to H_0 is

$$d_+ = \frac{|\langle \mathbf{w} \cdot \mathbf{x}^+ \rangle + b|}{\|\mathbf{w}\|} = \frac{|1|}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$$
[Eq.6]

- Similarly, the distance d_- from $\mathbf{x}^-$ to H_0 is

$$d_- = \frac{|\langle \mathbf{w} \cdot \mathbf{x}^- \rangle + b|}{\|\mathbf{w}\|} = \frac{|-1|}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$$
[Eq.7]

- As a result, the margin is:

$$margin = d_+ + d_- = \frac{2}{\|\mathbf{w}\|}$$
[Eq.8]

SVM: learning with max margin (1)

- *SVM learns a classifier H_0 with a maximum margin*, i.e., the hyperplane that has the greatest margin among all possible hyperplanes.

- This learning principle can be formulated as the following quadratic optimization problem:
  - *Find $\mathbf{w}$ and b that maximize*

$$margin = \frac{2}{\|\mathbf{w}\|}$$

  - *and satisfy the below conditions for any training instance $\mathbf{x_i}$:*

$$\begin{cases} \langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b \geq 1, \text{if } y_i = 1 \\ \langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b \leq -1, \text{if } y_i = -1 \end{cases}$$

SVM: learning with max margin (2)

- Learning SVM is equivalent to the following minimization problem:

  □ Minimize $\dfrac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2}$ [Eq.9]

  □ Conditioned on $\begin{cases} \langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b \geq 1, & if\ y_i = 1 \\ \langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b \leq -1, & if\ y_i = -1 \end{cases}$

- Note, it can be reformulated as:

  □ Minimize $\dfrac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2}$ [Eq.10]

  (P)

  □ Conditioned on $y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1, \quad \forall i = 1..r$

- This is a *constrained optimization problem*.

Constrained optimization (1)

- Consider the problem:

  Minimize *f(x)* conditioned on *g(x)* = 0

- *Necessary condition:* a solution $\mathbf{x}_0$ will satisfy

$$\begin{cases} \dfrac{\partial}{\partial \mathbf{x}}\left(f(\mathbf{x}) + \alpha g(\mathbf{x})\right)\Big|_{\mathbf{x}=\mathbf{x}_0} = 0 \\ g(\mathbf{x}) = 0 \end{cases};$$

  □ Where α is a *Lagrange multiplier*.

- In the cases of many constraints ($g_i(\mathbf{x})=0$ for i=1…r), a solution $\mathbf{x}_0$ will satisfy:

$$\begin{cases} \dfrac{\partial}{\partial \mathbf{x}}\left(f(\mathbf{x}) + \sum_{i=1}^{r} \alpha_i g_i(\mathbf{x})\right)\Big|_{\mathbf{x}=\mathbf{x}_0} = 0 \\ g_i(\mathbf{x}) = 0 \end{cases};$$

Constrained optimization (2)

- Consider the problem with inequality constraints:

    Minimize *f(**x**)* conditioned on $g_i(**x**) \leq 0$

- *Necessary condition:* a solution $\mathbf{x}_0$ will satisfy

$$\begin{cases} \dfrac{\partial}{\partial \mathbf{x}}\left(f(\mathbf{x}) + \sum_{i=1}^{r} \alpha_i g_i(\mathbf{x}) \right)\Bigg|_{\mathbf{x}=\mathbf{x_0}} = 0; \\ g_i(\mathbf{x}) \leq 0 \end{cases}$$

  □ Where $\alpha_i \geq 0$ is a Lagrange multiplier.

- $L = f(\mathbf{x}) + \sum_{i=1}^{r} \alpha_i g_i(\mathbf{x})$ is known as the Lagrange function.

  □ **x** is called *primal variable* (biến gốc)

  □ α is called *dual variable* (biến đối ngẫu)

SVM: learning with max margin (3)

- The Lagrange function for problem [Eq. 10] is

$$L(\mathbf{w}, b, \alpha) = \frac{1}{2}\langle \mathbf{w} \cdot \mathbf{w} \rangle - \sum_{i=1}^{r} \alpha_i [y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) - 1] \qquad \text{[Eq.11a]}$$

  □ Where each $\alpha_i \geq 0$ is a Lagrange multiplier.

- Solving [Eq. 10] is equivalent to the following minimax problem:

$$\min_{\mathbf{w},b} \max_{\alpha \geq 0} L(\mathbf{w}, b, \alpha) \qquad \text{[Eq.11b]}$$

$$= \min_{\mathbf{w},b} \max_{\alpha \geq 0} \left(\frac{1}{2}\langle \mathbf{w} \cdot \mathbf{w} \rangle - \sum_{i=1}^{r} \alpha_i [y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) - 1] \right)$$

SVM: learning with max margin (4)

- The *primal problem* [Eq. 10] can be derived by solving:

$$\max_{\alpha \geq 0} L(\mathbf{w}, b, \alpha) = \max_{\alpha \geq 0} \left(\frac{1}{2} \langle \mathbf{w} \cdot \mathbf{w} \rangle - \sum_{i=1}^{r} \alpha_i [y_i (\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) - 1] \right)$$

- Its *dual problem* (đối ngẫu) can be derived by solving:

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b, \alpha) = \min_{\mathbf{w}, b} \left(\frac{1}{2} \langle \mathbf{w} \cdot \mathbf{w} \rangle - \sum_{i=1}^{r} \alpha_i [y_i (\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) - 1] \right)$$

- It is known that the optimal solution to [Eq. 10] will satisfy some conditions which is called the **Karush-Kuhn-Tucker** (KKT) conditions.

SVM: Karush-Kuhn-Tucker

$$\frac{\partial L}{\partial \mathbf{w}} = \mathbf{w} - \sum_{i=1}^{r} \alpha_i y_i \mathbf{x_i} = 0 \qquad \text{[Eq.12]}$$

$$\frac{\partial L}{\partial b} = -\sum_{i=1}^{r} \alpha_i y_i = 0 \qquad \text{[Eq.13]}$$

$$y_i \left(\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b \right) - 1 \geq 0, \ \forall \mathbf{x_i} \ (i = 1..r) \qquad \text{[Eq.14]}$$

$$\alpha_i \geq 0 \qquad \text{[Eq.15]}$$

$$\alpha_i \left(y_i \left(\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b \right) - 1 \right) = 0 \qquad \text{[Eq.16]}$$

- The last equation [Eq. 16] comes from a nice result from the duality theory.

  □ Note: any $\alpha_i > 0$ will imply that the associated point $\mathbf{x_i}$ lies in a boundary hyperplane (H_+ or H_-).

  □ Such a boundary point is named as a *support vector*.

  □ A non-support vector will correspond to $\alpha_i = 0$.

SVM: learning with max margin (5)

- In general, the KKT conditions do not guarantee the optimality of the solution.

- Fortunately, due to the convexity of the primal problem [Eq.10], the *KKT conditions are both necessary and sufficient to assure the global optimality of the solution. It means a vector satisfying all KKT conditions provides the globally optimal classifier.*

  □ Convex optimization is 'easy' in the sense that we always can find a good solution with a provable guarantee.

  □ There are many algorithms in the literature, but most are iterative.

- In fact, problem [Eq.10] is pretty hard to derive an efficient algorithm. Therefore, its **dual problem** is more preferable.

SVM: the dual form (1)

- Remember that the dual counterpart of [Eq.10] is

$$\min_{\mathbf{w},b} L(\mathbf{w},b,\alpha) = \min_{\mathbf{w},b} \left(\frac{1}{2} \langle \mathbf{w} \cdot \mathbf{w} \rangle - \sum_{i=1}^{r} \alpha_i [y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) - 1] \right)$$

- By taking the gradient of L($\mathbf{w}$,b,α) in variables ($\mathbf{w}$,b) and zeroing it, we can find the following dual function:

$$L_D(\alpha) = \sum_{i=1}^{r} \alpha_i - \frac{1}{2} \sum_{i,j=1}^{r} \alpha_i \alpha_j y_i y_j \langle \mathbf{x}_i \cdot \mathbf{x}_j \rangle \qquad \text{[Eq.17]}$$

SVM: the dual form (2)

- Solving problem [Eq.10] is equivalent to solving its dual problem below:

  - Maximize $\quad L_D(\boldsymbol{\alpha}) = \sum_{i=1}^{r} \alpha_i - \frac{1}{2} \sum_{i,j=1}^{r} \alpha_i \alpha_j y_i y_j \langle \mathbf{x_i} \cdot \mathbf{x_j} \rangle \qquad$ [Eq.18]

    (D)

  - Such that $\qquad \begin{cases} \sum_{i=1}^{r} \alpha_i y_i = 0 \\ \alpha_i \geq 0, \forall i = 1..r \end{cases}$

- The constraints in **(D)** is much simpler than those of the primal problem. Therefore, deriving an efficient method to solve this problem might be easier.

  - However, existing algorithms for this problem are iterative and complicated. Therefore, we will not discuss any algorithm in detail !

SVM: the optimal classifier

- Once the dual problem is solved for $\boldsymbol{\alpha}$, we can recover the optimal solution to problem [Eq.10] by using the KKT.

- Let SV be the set of all support vectors

  - SV is a subset of the training data.

  - $\alpha_i > 0$ suggests that $\mathbf{x_i}$ is a support vector.

- We can compute $\mathbf{w}^*$ by using [Eq.12]. So:

  $$\mathbf{w}^* = \sum_{i=1}^{r} \alpha_i y_i \mathbf{x_i} = \sum_{\mathbf{x_i} \in SV} \alpha_i y_i \mathbf{x_i}; \qquad (\text{due to } \alpha_j = 0 \text{ for any } \mathbf{x_j} \text{ not in SV})$$

- To find b*, we take an index k such that $\alpha_k > 0$:

  - It means $y_k(\langle \mathbf{w}^* \cdot \mathbf{x}_k \rangle + b^*) - 1 = 0$ due to [Eq.16].

  - Hence,

  - $b^* = y_k - \langle \mathbf{w}^* \cdot \mathbf{x}_k \rangle$

SVM: classifying new instances

- The decision boundary is

$$f(\mathbf{x}) = \langle \mathbf{w}^* \cdot \mathbf{x} \rangle + b^* = \sum_{\mathbf{x_i} \in SV} \alpha_i y_i \langle \mathbf{x_i} \cdot \mathbf{x} \rangle + b^* = 0 \qquad \text{[Eq.19]}$$

- For a new instance **z**, we compute:

$$sign(\langle \mathbf{w}^* \cdot \mathbf{z} \rangle + b^*) = sign\left(\sum_{\mathbf{x_i} \in SV} \alpha_i y_i \langle \mathbf{x_i} \cdot \mathbf{z} \rangle + b^* \right) \qquad \text{[Eq.20]}$$

  □ If the result is 1, **z** will be assigned to the possitive class; otherwise **z** will be assigned to the negative class.

- Note that this classification principle

  □ Just depends on the support vectors.

 □ Just needs to compute some dot products.

---

2. Soft-margin SVM

- *What if the two classes are not linearly separable?*
  (Trường hợp 2 lớp không thể phân tách tuyến tính thì sao?)

  □ Linear separability is ideal in practice.

  □ Data are often noisy or erronous, making two classes overlapping (nhiễu/lỗi có thể làm 2 lớp giao nhau)

- In the case of linear separability:

  □ *Minimize* $\dfrac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2}$

  □ *Conditioned on* $y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1, \quad \forall i = 1..r$

- In the cases of **noises** or **overlapping**, those constraints may never meet simultaneously.

  □ It means we cannot solve for **w*** and b*.

Example of inseparability

- Noisy points $\mathbf{x}_a$ and $\mathbf{x}_b$ are mis-placed.

Relaxing the constraints

- To work with noises/errors, we need to relax the constraints about margin by using some slack variables ξ_i (≥ 0):
  (Ta sẽ mở rộng ràng buộc về lề bằng cách thêm biến bù)

  $$\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b \geq 1 - \xi_i \qquad \text{if } y_i = 1$$

  $$\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b \leq -1 + \xi_i \qquad \text{if } y_i = -1$$

  □ For a noisy/erronous point $\mathbf{x_i}$, we have: $\xi_i > 1$

  □ Otherwise $\xi_i = 0$.

- Therefore, we have the following constraints for the cases of linear inseparability:

  $$y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1 - \xi_i \qquad \text{for all } i = 1 \ldots r$$

  $$\xi_i \geq 0 \qquad \text{for all } i = 1 \ldots r$$

Penalty on noises/errors

- We should enclose some information on noises/errors into the objective function when learning
  (ta nên đính thêm thông tin về nhiễu/lỗi vào hàm mục tiêu)

  □ Otherwise, the resulting classifier easily underfits the data.

- A penalty term will be used so that learning is to minimize

$$\frac{\langle W, W \rangle}{2} + C \sum_{i=1}^{r} \xi_i^k$$

  □ Where C (>0) is the penalty constant (hằng số phạt).

  □ The greater C, the heavier the penalty on noises/errors.

- $k = 1$ is often used in practice, due to simplicity for solving the optimization problem.

The new optimization problem

- *Minimize*

$$\frac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2} + C \sum_{i=1}^{r} \xi_i$$

[Eq.21]

  □ *Conditioned on*

$$\begin{cases} y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1 - \xi_i, & \forall i = 1..r \\ \xi_i \geq 0, & \forall i = 1..r \end{cases}$$

- This problem is called **Soft-margin SVM**.

- It is equivalent to minimize the following function

$$\left[\frac{1}{r} \sum_{i=1}^{r} \max(0, 1 - y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b)) \right] + \lambda \|\mathbf{w}\|_2^2$$

  □ $\max(0, 1 - y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b))$ is called *Hinge loss*

  □ Some popular losses: squared error, cross entropy, hinge

  □ $\lambda > 0$ is a constant

The new optimization problem

■ Its Lagrange function is

$$L = \frac{1}{2}\langle \mathbf{w} \cdot \mathbf{w} \rangle + C\sum_{i=1}^{r}\xi_i - \sum_{i=1}^{r}\alpha_i[y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) - 1 + \xi_i] - \sum_{i=1}^{r}\mu_i\xi_i$$

[Eq.22]

□ Where α_i (≥0) and μ_i (≥0) are Lagrange multipliers.

Karush-Kuhn-Tucker conditions (1)

$$\frac{\partial L_P}{\partial \mathbf{w}} = \mathbf{w} - \sum_{i=1}^{r}\alpha_i y_i \mathbf{x_i} = 0$$

[Eq.23]

$$\frac{\partial L_P}{\partial b} = -\sum_{i=1}^{r}\alpha_i y_i = 0$$

[Eq.24]

$$\frac{\partial L_P}{\partial \xi_i} = C - \alpha_i - \mu_i = 0, \quad \forall i = 1..r$$

[Eq.25]

Karush-Kuhn-Tucker conditions (2)

$$y_i\left(\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b\right) - 1 + \xi_i \geq 0, \quad \forall i = 1..r \qquad \text{[Eq.26]}$$

$$\xi_i \geq 0 \qquad \text{[Eq.27]}$$

$$\alpha_i \geq 0 \qquad \text{[Eq.28]}$$

$$\mu_i \geq 0 \qquad \text{[Eq.29]}$$

$$\alpha_i\left(y_i\left(\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b\right) - 1 + \xi_i\right) = 0 \qquad \text{[Eq.30]}$$

$$\mu_i \xi_i = 0 \qquad \text{[Eq.31]}$$

The dual problem

- *Maximize*

$$L_D(\boldsymbol{\alpha}) = \sum_{i=1}^{r} \alpha_i - \frac{1}{2} \sum_{i,j=1}^{r} \alpha_i \alpha_j y_i y_j \langle \mathbf{x_i} \cdot \mathbf{x_j} \rangle$$

  □ *Such that*

$$\begin{cases} \sum_{i=1}^{r} \alpha_i y_i = 0 \\ 0 \leq \alpha_i \leq C, \quad \forall i = 1..r \end{cases}$$

  [Eq.32]

- Note that neither ξ nor μ_i appears in the dual problem.
- This problem is almost similar with that [Eq.18] in the case of linearly separable classification.
- The only difference is the constraint: $\alpha_i \leq C$

Soft-margin SVM: the optimal classifier

- Once the dual problem is solved for α, we can recover the optimal solution to problem [Eq.21].

- Let SV be the set of all *support/noisy vectors*

  - □ SV is a subset of the training data.

  - □ $\alpha_i > 0$ suggests that $\mathbf{x_i}$ is a support/noisy vector.

- We can compute $\mathbf{w}^*$ by using [Eq.12]. So:

$$\mathbf{w}^* = \sum_{i=1}^{r} \alpha_i y_i \mathbf{x_i} = \sum_{\mathbf{x_i} \in SV} \alpha_i y_i \mathbf{x_i}; \qquad \text{(due to } \alpha_j = 0 \text{ for any } \mathbf{x_j} \text{ not in SV)}$$

- To find b*, we take an index k such that $C > \alpha_k > 0$:

  - □ It means $\xi_k = 0$ due to [Eq.25] and [Eq.31];

  - □ And $y_k(\langle \mathbf{w}^* \cdot \mathbf{x}_k \rangle + b^*)$ -1 = 0 due to [Eq.30].

  - □ Hence, b* = y_k - $\langle \mathbf{w}^* \cdot \mathbf{x}_k \rangle$

Some notes

- From [Eq.25-31] we conclude that

$$
\begin{array}{llll}
\text{If} & \alpha_i = 0 & \text{then} & y_i(\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b) \geq 1, & \text{and } \xi_i = 0 \\
\text{If} & 0 < \alpha_i < C & \text{then} & y_i(\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b) = 1, & \text{and } \xi_i = 0 \\
\text{If} & \alpha_i = C & \text{then} & y_i(\langle \mathbf{w} \cdot \mathbf{x_i} \rangle + b) < 1, & \text{and } \xi_i > 0
\end{array}
$$

- The classifier can be expressed as a *linear combination* of few training points.

  - □ Most training points lie outside the margin area: $\alpha_i = 0$

  - □ The support vectors lie in the marginal hyperplanes: $0 < \alpha_i < C$

  - □ The noisy/erronous points will associate with $\alpha_i = C$

- Hence the optimal classifier is a very sparse combination of the training data.

Soft-margin SVM: classifying new instances

- The decision boundary is

$$f(\mathbf{x}) = \langle \mathbf{w}^* \cdot \mathbf{x} \rangle + b^* = \sum_{\mathbf{x_i} \in SV} \alpha_i y_i \langle \mathbf{x_i} \cdot \mathbf{x} \rangle + b^* = 0 \qquad \text{[Eq.19]}$$

- For a new instance **z**, we compute:

$$sign(\langle \mathbf{w}^* \cdot \mathbf{z} \rangle + b^*) = sign\left(\sum_{\mathbf{x_i} \in SV} \alpha_i y_i \langle \mathbf{x_i} \cdot \mathbf{z} \rangle + b^* \right) \qquad \text{[Eq.20]}$$

  □ If the result is 1, **z** will be assigned to the possitive class; otherwise **z** will be assigned to the negative class.

- *Note: it is important to choose a good value of C, since it significantly affects performance of SVM.*

  □ We often use a validation set to choose a value for C.

Linear SVM: summary

- Classification is based on a separating hyperplane.
- Such a hyperplane is represented as a combination of some support vectors.
- The determination of support vectors reduces to solve a quadratic programming problem.
- In the dual problem and the separating hyperplane, dot products can be used in place of the original training data.

 □ This is the door for us to learn a nonlinear classifier.

3. Non-linear SVM

- Consider the case in which our data are not linearly separable
  - □ This may often happen in practice
- How about using a non-linear function?
- **Idea of Non-linear SVM:**
  - □ **Step 1:** transform the input into another space, which often has *higher dimensions*, so that the projection of data is linearly separable
  - □ **Step 2:** use linear SVM in the new space

Non-linear SVM

- **Input space:** initial representation of data
- **Feature space:** the new space after the transformation

Input space $\xrightarrow{}$ Feature space

$$\phi(x)$$
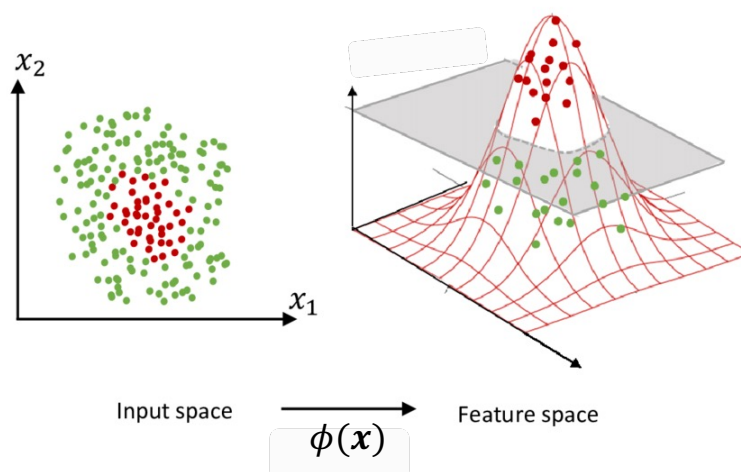
Non-linear SVM: transformation

- Our idea is to map the input **x** to a new representation, using a non-linear mapping

$$\phi: \ X \longrightarrow F$$
$$x \longmapsto \phi(x)$$

- In the feature space, the original training data $\{(x_1, y_1), (x_2, y_2), \dots, (x_r, y_r)\}$ are represented by

$$\{(\phi(x_1), y_1), (\phi(x_2), y_2), \dots, (\phi(x_r), y_r)\}$$

Non-linear SVM: transformation

- Consider the input space to be 2-dimensional, and we choose the following map

$$\phi: \ X \longrightarrow F$$
$$(x_1, x_2) \longmapsto (x_1^2, x_2^2, \sqrt{2}x_1x_2)$$

- So instance **x** = (2, 3) will have the representation in the feature space as

$$\phi(\mathbf{x}) = (4, 9, 8.49)$$

Non-linear SVM: learning & prediction

- **Training problem:**

  Minimize
  $$L_P = \frac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2} + C \sum_{i=1}^{r} \xi_i \qquad \text{[Eq.34]}$$

  Such that
  $$\begin{cases} y_i \left(\langle \mathbf{w} \cdot \phi(\mathbf{x_i}) \rangle + b \right) \geq 1 - \xi_i, \quad \forall i = 1..r \\ \xi_i \geq 0, \quad \forall i = 1..r \end{cases}$$

- **The dual problem:**

  Maximize
  $$L_D = \sum_{i=1}^{r} \alpha_i - \frac{1}{2} \sum_{i,j=1}^{r} \alpha_i \alpha_j y_i y_j \langle \phi(\mathbf{x_i}) \cdot \phi(\mathbf{x_j}) \rangle \qquad \text{[Eq.35]}$$

  Such that
  $$\begin{cases} \sum_{i=1}^{r} \alpha_i y_i = 0 \\ 0 \leq \alpha_i \leq C, \quad \forall i = 1..r \end{cases}$$

- **Classifier:**

  $$f(\mathbf{z}) = \langle \mathbf{w}^*, \phi(\mathbf{z}) \rangle + b^* = \sum_{\mathbf{x_i} \in SV} \alpha_i y_i \langle \phi(\mathbf{x_i}), \phi(\mathbf{z}) \rangle + b^* \qquad \text{[Eq.36]}$$

Non-linear SVM: difficulties

- How to find the mapping?
  - An intractable problem
- The curse of dimensionality
  - As the dimensionality increases, the volume of the space increases so fast that the available data become sparse.
  - This sparsity is problematic.
  - Increasing the dimensionality will require significantly more training data.

  Dữ liệu dù thu thập được lớn đến đâu thì cũng là **quá nhỏ** so với không gian của chúng

Non-linear SVM: Kernel functions

- An explicit form of a tranformation is not necessary

- **The dual problem:**

  Maximize
  $$L_D = \sum_{i=1}^{r} \alpha_i - \frac{1}{2} \sum_{i,j=1}^{r} \alpha_i \alpha_j y_i y_j \langle \phi(\mathbf{x_i}) \cdot \phi(\mathbf{x_j}) \rangle$$

  Such that
  $$\begin{cases} \sum_{i=1}^{r} \alpha_i y_i = 0 \\ 0 \leq \alpha_i \leq C, \quad \forall i = 1..r \end{cases}$$

- **Classifier:** $f(\mathbf{z}) = \langle \mathbf{w}^*, \phi(\mathbf{z}) \rangle + b^* = \sum_{x_i \in SV} \alpha_i y_i \langle \phi(x_i), \phi(\mathbf{z}) \rangle + b^*$

- Both require only the inner product $\langle \phi(\mathbf{x}), \phi(\mathbf{z}) \rangle$

- **Kernel trick:** Nonlinear SVM can be used by replacing those inner products by evaluations of some *kernel function*

  $$K(\mathbf{x,z}) = \langle \phi(\mathbf{x}), \phi(\mathbf{z}) \rangle \qquad \text{[Eq.37]}$$

Kernel functions: example

- Polynomial
  $$K(\boldsymbol{x}, \boldsymbol{z}) = \langle \boldsymbol{x}, \boldsymbol{z} \rangle^d$$

- Consider the polynomial with degree d=2. For any vectors $\mathbf{x}=(x_1,x_2)$ and $\mathbf{z}=(z_1,z_2)$

$$\begin{aligned} \langle \boldsymbol{x}, \boldsymbol{z} \rangle^2 \quad &= \quad (x_1 z_1 + x_2 z_2)^2 \\ &= \quad x_1^2 z_1^2 + 2 x_1 z_1 x_2 z_2 + x_2^2 z_2^2 \\ &= \quad \langle (x_1^2, x_2^2, \sqrt{2} x_1 x_2), (z_1^2, z_2^2, \sqrt{2} z_1 z_2) \rangle \\ &= \quad \langle \phi(\boldsymbol{x}), \phi(\boldsymbol{z}) \rangle = K(\boldsymbol{x}, \boldsymbol{z}) \end{aligned}$$

  □ Where $\phi(\boldsymbol{x}) = (x_1^2, x_2^2, \sqrt{2} x_1 x_2)$.

- Therefore the polynomial is the product of two vectors $\phi(\boldsymbol{x})$ and $\phi(\boldsymbol{z})$.

Kernel functions: popular choices

- Polynomial

$$K(\mathbf{x},\mathbf{z}) = \left(\langle \mathbf{x} \cdot \mathbf{z} \rangle + \theta\right)^d; \ \text{trong đó}: \ \theta \in R, d \in N$$

- Gaussian radial basis function (RBF)

$$K(\mathbf{x},\mathbf{z}) = e^{-\frac{\|\mathbf{x}-\mathbf{z}\|^2}{2\sigma}}; \ \text{trong đó}: \ \sigma > 0$$

- Sigmoid

$$K(x, z) = \tanh(\gamma \langle x \cdot z \rangle + r), \ \text{trong đó} \ \gamma, r \in R$$

- What conditions ensure a kernel function?
  Mercer's theorem

SVM: summary

- SVM works with real-value attributes
  - □ Any nominal attribute need to be transformed into a real one
- The learning formulation of SVM focuses on 2 classes
  - □ How about a classification problem with > 2 classes?
  - □ One-vs-the-rest, one-vs-one: a multiclass problem can be solved by reducing to many different problems with 2 classes
- The decision function is simple, but may be hard to interpret
  - □ It is more serious if we use some kernel functions

SVM: some packages

- LibSVM:
  - http://www.csie.ntu.edu.tw/~cjlin/libsvm/

- Linear SVM for large datasets:
  - http://www.csie.ntu.edu.tw/~cjlin/liblinear/
  - http://www.cs.cornell.edu/people/tj/svm_light/svm_perf.html

- Scikit-learn in python:
  - http://scikit-learn.org/stable/modules/svm.html

- SVM^{light}:
  - http://www.cs.cornell.edu/people/tj/svm_light/index.html

References

- B. Liu. Web Data Mining: Exploring Hyperlinks, Contents, and Usage Data. Springer, 2006.
- C. J. C. Burges. A Tutorial on Support Vector Machines for Pattern Recognition. Data Mining and Knowledge Discovery, 2(2): 121-167, 1998.
- Cortes, Corinna, and Vladimir Vapnik. "Support-vector networks." *Machine learning* 20.3 (1995): 273-297.

Exercises

- What is the main difference between SVM and KNN?

- How many support vectors are there in the worst case? Why?

- The meaning of the constant C in SVM? Compare the role of C in SVM with that of λ in Ridge regression.