



Machine Learning

(Học máy – IT3190E)

Khoat Than

School of Information and Communication Technology

Hanoi University of Science and Technology

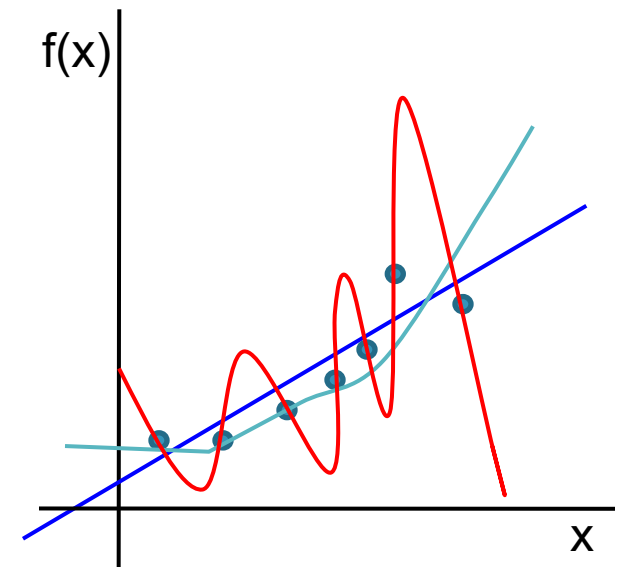
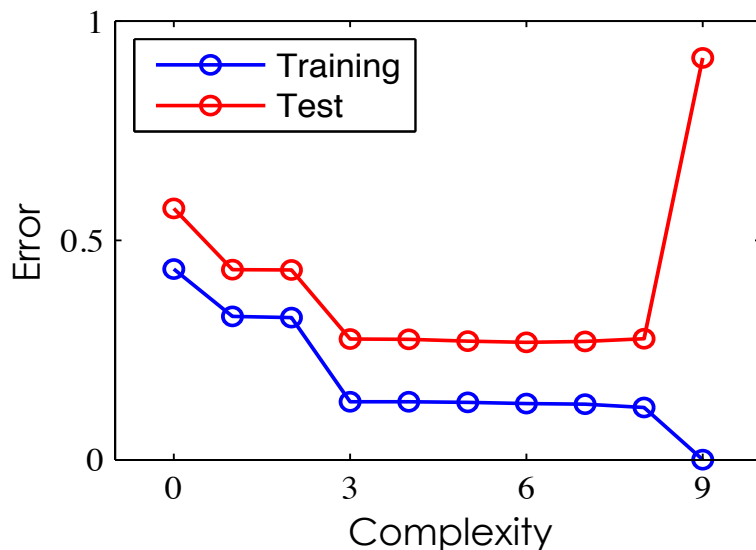
2022

Contents

- Introduction to Machine Learning
- Supervised learning
- Probabilistic modeling
- **Regularization**
- Reinforcement learning
- Practical advice

Revisiting overfitting

- The complexity of the learned function: $y = \hat{f}(x, \mathbf{D})$
 - For a given training data \mathbf{D} : the more complicated \hat{f} , the more possibility that \hat{f} fits \mathbf{D} better.
 - For a given \mathbf{D} : there exist many functions that fit \mathbf{D} perfectly (i.e., no error on \mathbf{D}).
 - However, those functions might generalize badly.



The Bias-Variance Decomposition

- Consider $y = f(x) + \epsilon$ as the regression function
 - ❖ where $\epsilon \sim \mathcal{N}(0, \sigma^2)$ is a Gaussian noise with mean 0 and variance σ^2 .
 - ❖ ϵ may represent the *noise* due to measurement or data collection.
- Let $\hat{f}(x; \mathbf{D})$ be the regressor learned from a training data \mathbf{D}
- Note:
 - ❖ We want that $\hat{f}(x; \mathbf{D})$ approximates the truth $f(x)$ well.
 - ❖ $\hat{f}(x; \mathbf{D})$ is random, according to the randomness when collecting \mathbf{D} .
- For any x , the error made by $\hat{f}(x; \mathbf{D})$ is

$$\mathbb{E}_{D, \epsilon} \left(y(x) - \hat{f}(x; \mathbf{D}) \right)^2 = \sigma^2 + \text{Bias}^2 \left(\hat{f}(x; \mathbf{D}) \right) + \text{Var} \left(\hat{f}(x; \mathbf{D}) \right)$$

$$\text{Bias} \left(\hat{f}(x; \mathbf{D}) \right) = \mathbb{E}_D [f(x) - \hat{f}(x; \mathbf{D})]$$

$$\text{Var} \left(\hat{f}(x; \mathbf{D}) \right) = \mathbb{E}_D \left(\hat{f}(x; \mathbf{D}) - \mathbb{E}_D \hat{f}(x; \mathbf{D}) \right)^2$$

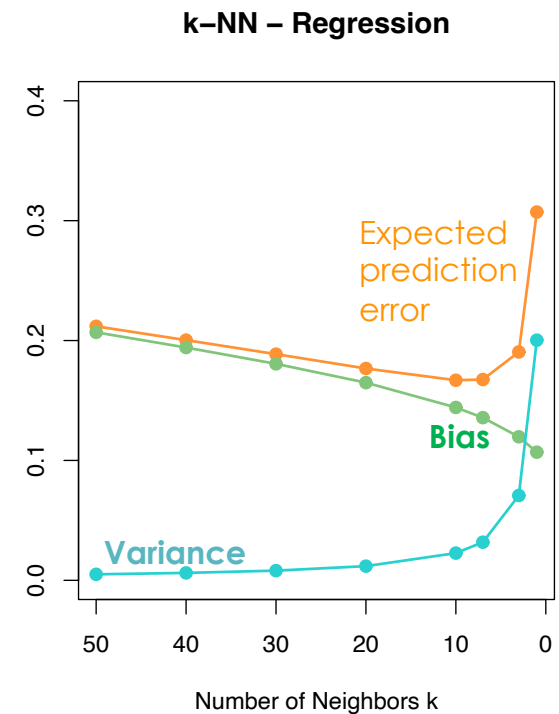
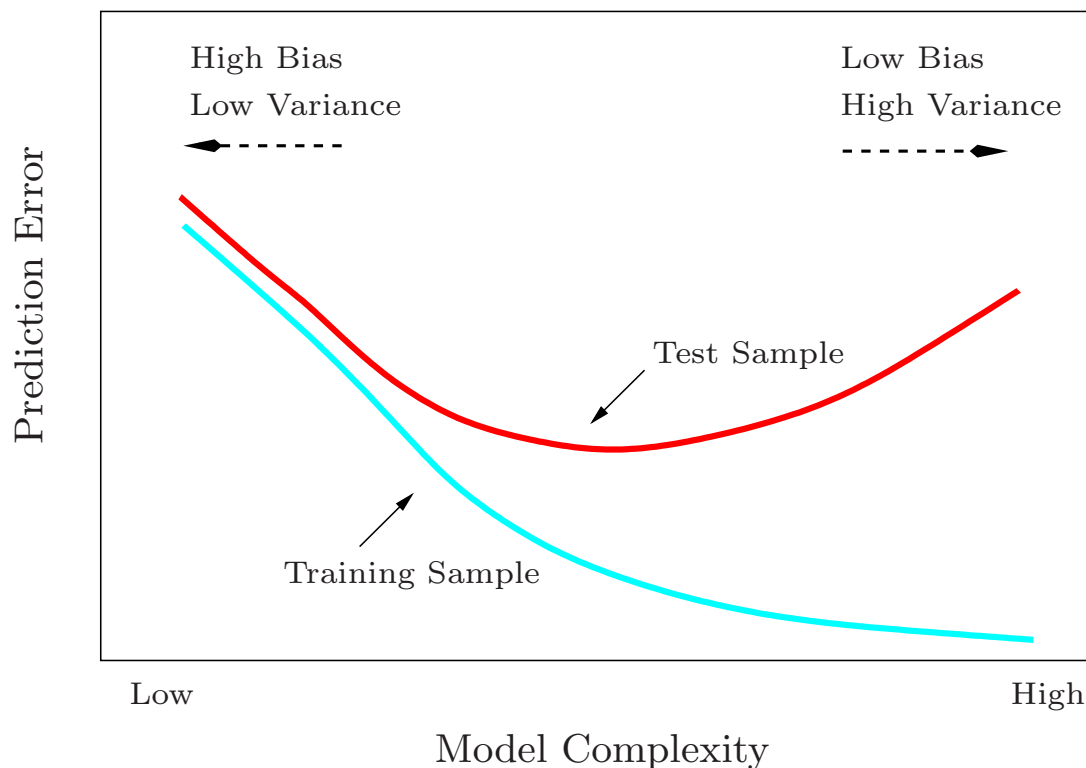
The Bias-Variance Decomposition (2)

$$\begin{aligned} \text{Error}(x) &= \sigma^2 + \text{Bias}^2 \left(\hat{f}(x; \mathbf{D}) \right) + \text{Var} \left(\hat{f}(x; \mathbf{D}) \right) \\ &= \text{Irreducible Error} + \text{Bias}^2 + \text{Variance} \end{aligned}$$

- This is known as the **Bias-Variance Decomposition**
 - ❖ *Irreducible Error*: cannot be avoided due to noises or uncontrolled factors
 - ❖ *Bias*: the average of our estimate differs from the true mean
 - ❖ *Variance*: the expected squared deviation of $\hat{f}(x; \mathbf{D})$ around its mean

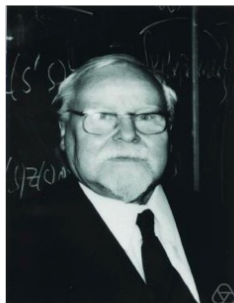
Bias-Variance tradeoff: classical view

- The more complex the model $\hat{f}(x; \mathbf{D})$ is, the more data points it can capture, and the lower the bias can be.
- ❖ However, higher complexity will make the model "move" more to capture the data points, and hence its variance will be larger.

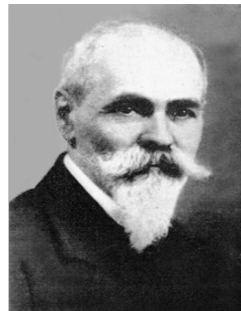


Regularization: introduction

- *Regularization* is now a popular and useful technique in ML.
- It is a technique to exploit further information to
 - Reduce overfitting in ML.
 - Solve ill-posed problems in Maths.
- The further information is often enclosed in a *penalty on the complexity* of $\hat{f}(x, \mathbf{D})$.
 - More penalty will be imposed on complex functions.
 - We prefer simpler functions among all that fit well the training data.



Tikhonov,
smoothing an ill-
posed problem



Zaremba, model
complexity
minimization



Bayes: priors
over parameters



Andrew Ng: need no
maths, but it prevents
overfitting!

Regularization in Ridge regression

- Learning a linear regressor by ordinary least squares (OLS) from a training data $\mathbf{D} = \{(x_1, y_1), \dots, (x_M, y_M)\}$ is reduced to the following problem:

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_2^2 = \arg \min_w \sum_{(x_i, y_i) \in \mathbf{D}} (y_i - w^T x_i)^2$$

- For Ridge regression, learning is reduced to

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_2^2$$

- Where λ is a positive constant.
- The term $\lambda \|w\|_2^2$ plays the role as *limiting the size/complexity of w* .
- λ allows us to trade off between fitness on \mathbf{D} and generalization on future observations.
- Ridge regression is a regularized version of OLS.

Regularization: the principle

- We need to learn a function $f(x, w)$ from the training set \mathbf{D}
 - x is a data example and belongs to **input space**.
 - w is the parameter and often belongs to a **parameter space \mathbf{W}** .
 - $\mathbf{F} = \{f(x, w) : w \in \mathbf{W}\}$ is the **function space**, parameterized by w .
- For many ML models, the training problem is often reduced to the following optimization:

$$w^* = \arg \min_{w \in \mathbf{W}} L(f(x, w), \mathbf{D}) \quad (1)$$

- w sometimes tells the size/complexity of that function.
 - $L(f(x, w), \mathbf{D})$ is an **empirical loss/risk** which depends on \mathbf{D} . This loss shows how well function f fits \mathbf{D} .
- Another view: $f^* = \arg \min_{f \in \mathbf{F}} L(f(x, w), \mathbf{D})$

Regularization: the principle

- Adding a penalty to (1), we consider

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) + \lambda g(w) \quad (2)$$

- Where $\lambda > 0$ is called *the regularization/penalty constant*.
- $g(w)$ measures the complexity of w . ($g(w) \geq 0$)
- $L(f(x, w), \mathbf{D})$ measures the goodness of function f on \mathbf{D} .
- The penalty (regularization) term: $\lambda g(w)$
 - Allows to trade off the fitness on \mathbf{D} and the generalization.
 - The greater λ , the heavier penalty, implying that $g(w)$ should be smaller.
 - In practice, λ should be neither too small nor too large.
(λ không nên quá lớn hoặc quá bé trong thực tế)

Regularization: popular types

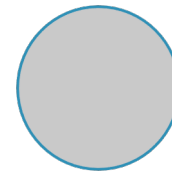
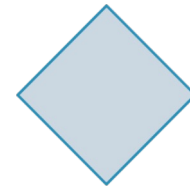
- $g(w)$ often relates to some norms when w is an n -dimensional vector.

□ L_0 -norm: $\|w\|_0$ counts the number of non-zeros in w .

□ L_1 -norm:
$$\|w\|_1 = \sum_{i=1}^n |w_i|$$

□ L_2 -norm:
$$\|w\|_2^2 = \sum_{i=1}^n w_i^2$$

□ L_p -norm:
$$\|w\|_p = \sqrt[p]{|w_1|^p + \dots + |w_n|^p}$$



Regularization in Ridge regression

- Ridge regression can be derived from OLS by adding a penalty term into the objective function when learning.
- Learning a regressor in Ridge is reduced to

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_2^2$$

- Where λ is a positive constant.
- The term $\lambda \|w\|_2^2$ plays the role as regularization.
- Large λ reduces the size of w .

Regularization in Lasso

- Lasso [Tibshirani, 1996] is a variant of OLS for linear regression by using L_1 to do regularization.
- Learning a linear regressor is reduced to

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_1$$

- Where λ is a positive constant.
- $\lambda \|w\|_1$ is the regularization term. Large λ reduces the size of w .
- Regularization here amounts to imposing a Laplace distribution (as prior) over each w_i , with density function:

$$p(w_i | \lambda) = \frac{\lambda}{2} e^{-\lambda |w_i|}$$

- The larger λ , the more possibility that $w_i = 0$.

Regularization in SVM

- Learning a classifier in SVM is reduced to the following problem:

- Minimize
$$\frac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2}$$

- Conditioned on $y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1, \quad \forall i = 1..r$

- In the cases of noises/errors, learning is reduced to

- Minimize
$$\frac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2} + C \sum_{i=1}^r \xi_i$$

- Conditioned on
$$\begin{cases} y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1 - \xi_i, & \forall i = 1..r \\ \xi_i \geq 0, & \forall i = 1..r \end{cases}$$

- $C(\xi_1 + \dots + \xi_r)$ is *the regularization term*.

Some other regularizations

- **Dropout:** (by Hilton and his colleagues, 2012)

- At each iteration of the training process, randomly drop out some parts and just update the other parts of our model.

- **Batch normalization** [Ioffe & Szegedy, 2015]

- Normalize the inputs at each neuron of a neural network
- Reduce input variance, easier training, faster convergence

- **Data augmentation**

- Produce different versions of an example in the training set, by adding simple noises, translation, rotation, cropping, ...
- Those versions are added to the training data set

- **Early stopping**

- Stop training early to avoid overtraining & reduce overfitting

Regularization: MAP role

- Under some conditions, we can view regularization as

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) + \lambda g(w)$$

Likelihood

Prior

- Where \mathbf{D} is a sample from a probability distribution whose log likelihood is $-L(f(x, w), \mathbf{D})$.

- w is a random variable and follows the prior with density
 $p(w) \propto \exp(-\lambda g(w))$

- Then $w^* = \arg \max_{w \in \mathcal{W}} \{-L(f(x, w), \mathbf{D}) - \lambda g(w)\}$

$$w^* = \arg \max_{w \in \mathcal{W}} \log \Pr(\mathbf{D}|w) + \log \Pr(w) = \arg \max_{w \in \mathcal{W}} \log \Pr(w|\mathbf{D})$$

- As a result, regularization in fact helps us to learn an MAP solution w^* .

Regularization: MAP in Ridge

- Consider the Gaussian regression model:

- w follows a Gaussian prior: $N(w \mid 0, \sigma^2 \rho^2)$.
- Variable $f = y - w^T x$ follows the Gaussian distribution $N(f \mid 0, \rho^2, w)$ with mean 0 and variance ρ^2 , and conditioned on w .

- Then the MAP estimation of f from the training data \mathbf{D} is

$$w^* = \operatorname{argmax}_w \log \Pr(w \mid D) = \operatorname{argmax}_w \log [\Pr(D \mid w) * \Pr(w)]$$

$$= \operatorname{argmin}_w \sum_{(x_i, y_i)} \frac{1}{2\rho^2} (y_i - w^T x_i)^2 + \frac{1}{2\sigma^2 \rho^2} w^T w - \text{constant}$$

$$= \operatorname{argmin}_w \sum_{(x_i, y_i)} (y_i - w^T x_i)^2 + \frac{1}{\sigma^2} w^T w$$

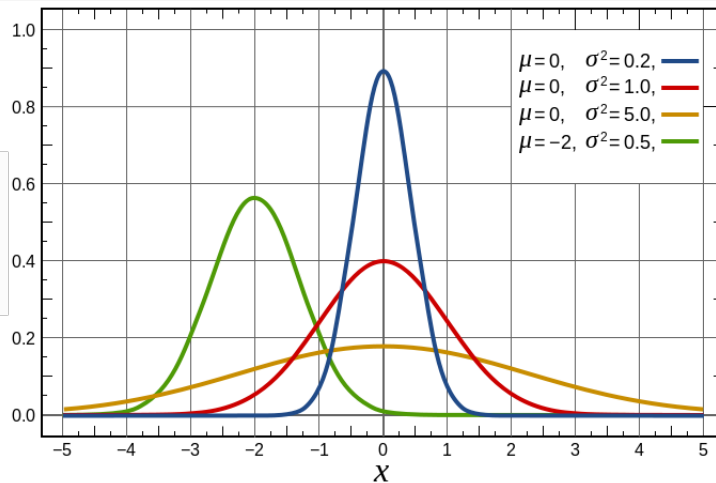
Ridge regression
@@

- Regularization using L_2 with penalty constant $\lambda = \sigma^2$.

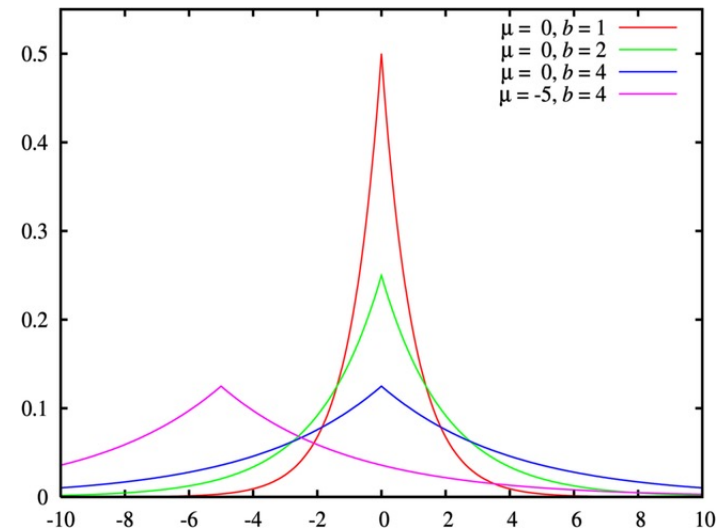
Regularization: MAP in Ridge & Lasso

- The regularization constant in Ridge: $\lambda = \sigma^{-2}$
- The regularization constant in Lasso: $\lambda = b^{-1}$
- Gaussian (left) and Laplace distribution (right)

$$f(x, \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



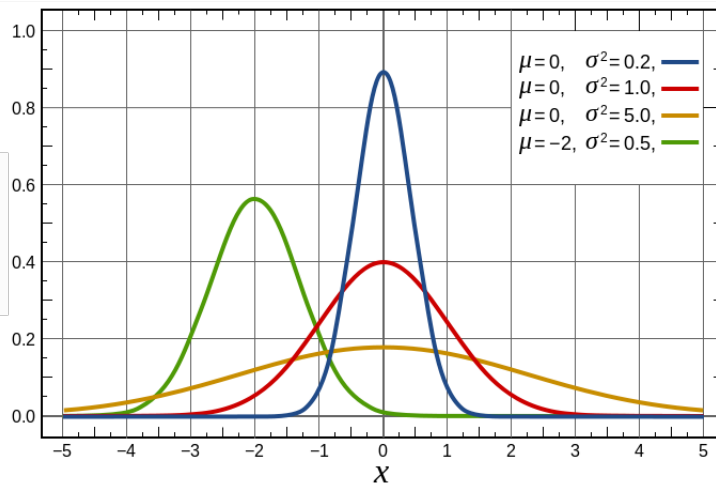
$$f(x|\mu, b) = \frac{1}{2b} \exp\left(-\frac{|x - \mu|}{b}\right)$$



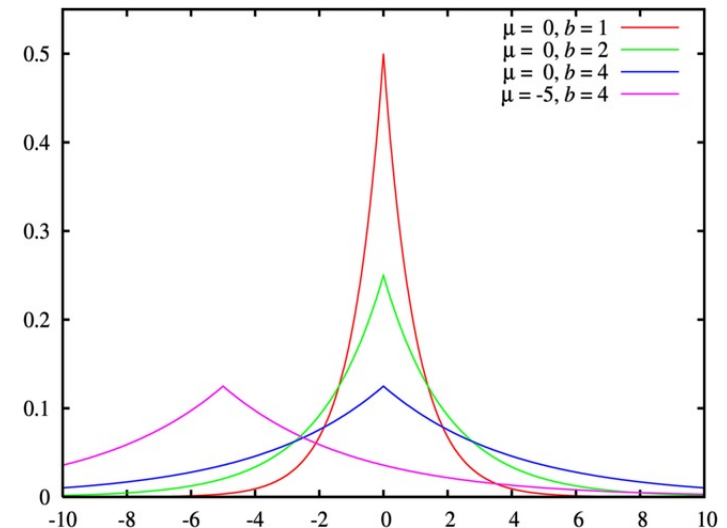
Regularization: limiting the search space

- The regularization constant in Ridge: $\lambda = \sigma^{-2}$
- The regularization constant in Lasso: $\lambda = b^{-1}$
- *The larger λ , the higher probability that x occurs around 0.*

$$f(x, \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



$$f(x|\mu, b) = \frac{1}{2b} \exp\left(-\frac{|x - \mu|}{b}\right)$$



Regularization: limiting the search space

- The regularized problem:

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) + \lambda g(w) \quad (2)$$

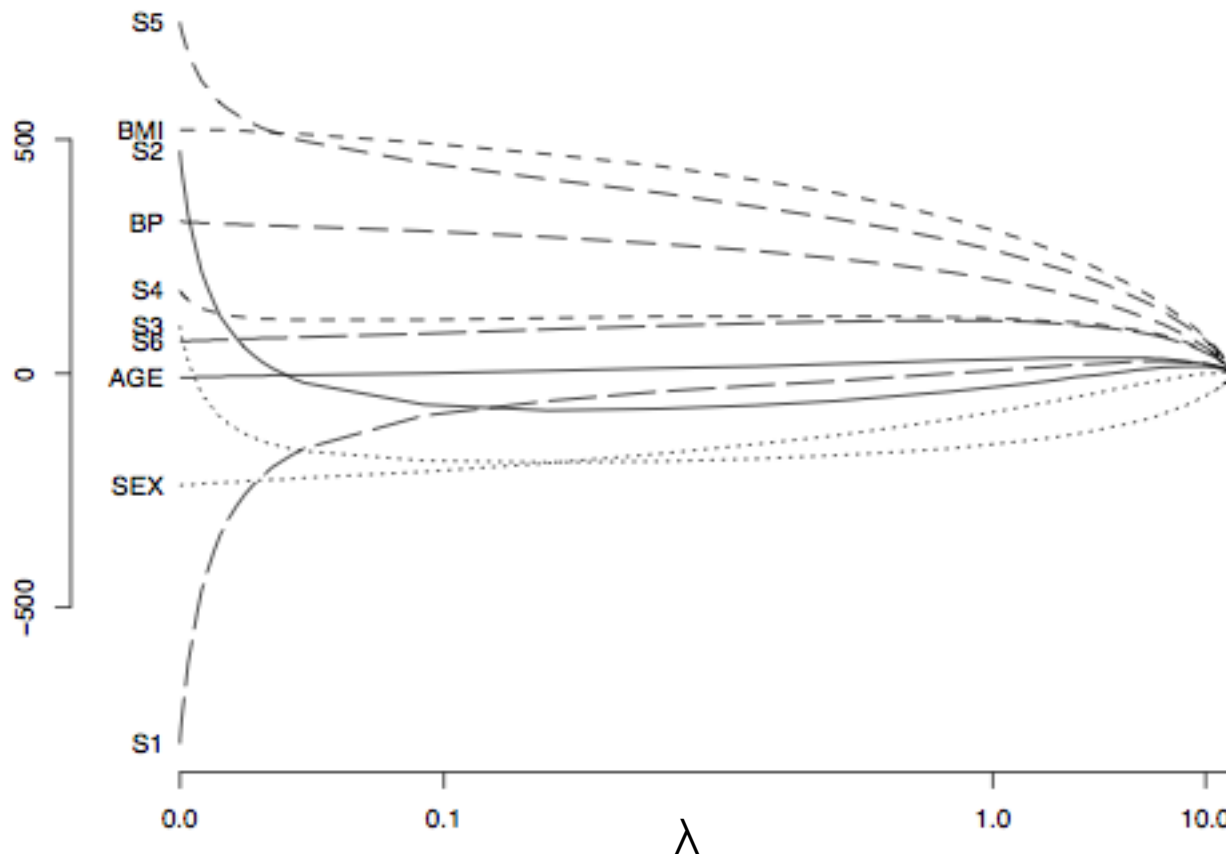
- A result from the optimization literature shows that (2) is equivalent to the following:

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) \quad \text{such that} \quad g(w) \leq s \quad (3)$$

- For some constant s .
- *Note that the constraint of $g(w) \leq s$ plays the role as limiting the search space of w .*

Regularization: effects of λ

- Vector $\mathbf{w}^* = (w_0, s1, s2, s3, s4, s5, s6, \text{Age}, \text{Sex}, \text{BMI}, \text{BP})$ changes when λ changes in Ridge regression.
 - \mathbf{w}^* goes to 0 as λ increases.



- Ridge regression was under investigation on a prostate dataset with 67 observations.
 - Performance was measured by RMSE (root mean square errors) and Correlation coefficient.

λ	0.1	1	10	100	1000	10000
RMSE	0.74	0.74	0.74	0.84	1.08	1.16
Correlation coefficient	0.77	0.77	0.78	0.76	0.74	0.73

- Too high or too low values of λ often result in bad predictions.
- Why??

Bias-Variance tradeoff: revisit

■ Classical view:

More complex model $\hat{f}(x; \mathbf{D})$

- ❖ Lower bias, higher variance

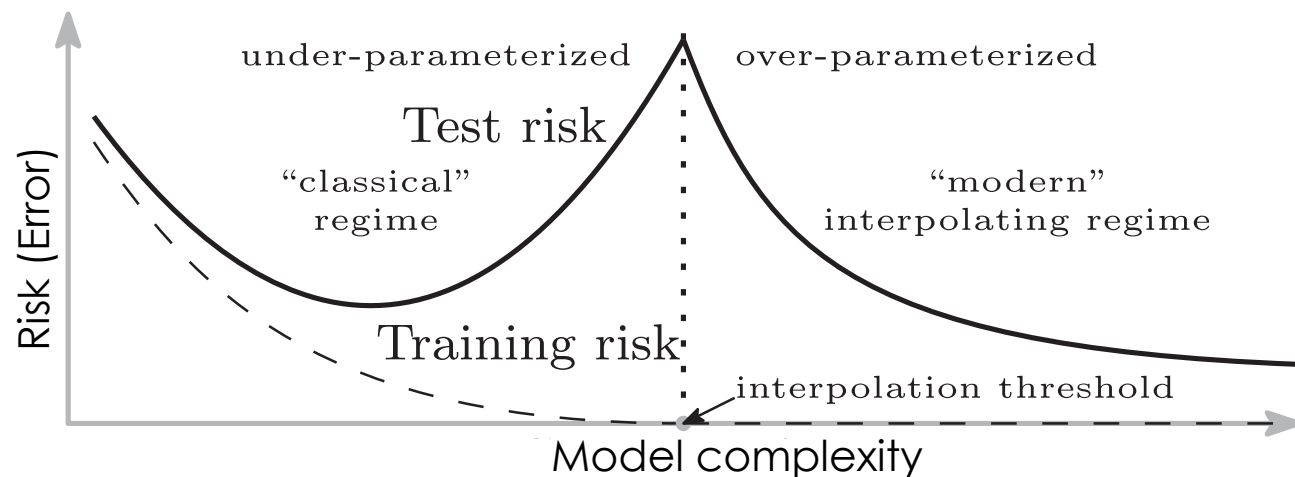
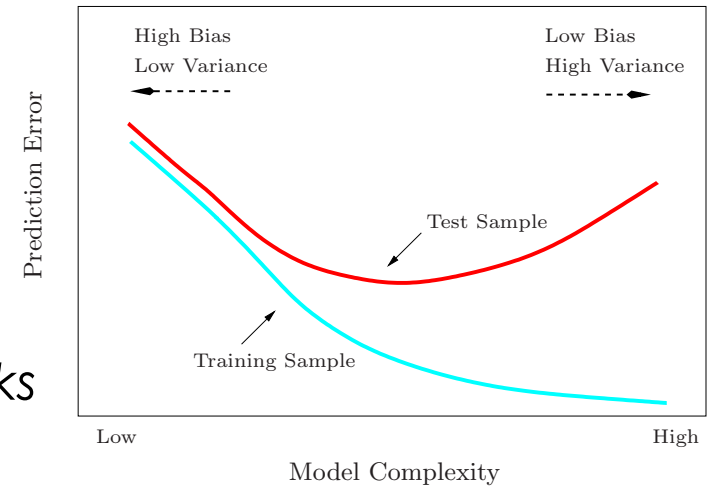
■ Modern phenomenon:

- ❖ Very rich models such as neural networks are trained to **exactly fit** the data, but often obtain **high accuracy** on test data [Belkin et al., 2019; Zhang et al., 2021]

- ❖ $\text{Bias} \cong 0$

- ❖ GPT-3, ResNets, VGG, StyleGAN, DALLE-3, ...

■ Why???



Regularization: summary

■ Advantages:

- Avoid overfitting.
- Limit the search space of the function to be learned.
- Reduce bad effects from noises or errors in observations.
- Might model data better. As an example, L_1 often work well with data/model which are inherently sparse.

■ Limitations:

- Consume time to select a good regularization constant.
- Might pose some difficulties to design an efficient algorithm.

References

- Belkin, M., Hsu, D., Ma, S., & Mandal, S. (2019). Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32), 15849-15854.
- Ioffe, S., & Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International Conference on Machine Learning* (pp. 448-456).
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. *Advances in Neural Information Processing Systems*, 25, 1097-1105.
- Tibshirani, R. (1996). *Regression shrinkage and selection via the Lasso*. *Journal of the Royal Statistical Society*, vol. 58(1), pp. 267-288.
- Trevor Hastie, Robert Tibshirani, Jerome Friedman. *The Elements of Statistical Learning*. Springer, 2009.
- Zhang, C., Bengio, S., Hardt, M., Recht, B., & Vinyals, O. (2021). Understanding deep learning (still) requires rethinking generalization. *Communications of the ACM*, 64(3), 107-115.